Development Report of
R Financial Option Pricing Package
`rmcop`

STAT0035
Student No. 20019563
Candidate No. WVTW5
Word count: 10395

April 22, 2023

Abstract

This report discusses the development of the R package `rmcop`, a fully R-based package for pricing financial options (see Section 1.3 for definitions). Pricing financial options is an essential fragment of financial engineering and quantitative finance. Because options consider the future in time, determining their fair price involves modelling uncertainty, so statistical methods are widely applied. Modern option pricing methods can be complex and product-specific, but most works are still built upon the basic frameworks of the Black-Scholes formula and Monte Carlo simulations. The package `rmcop` covers the R implementation of some of these fundamental option price models on various option types and styles and can be a good reference for constructing more complex R option pricing algorithms. The report also contains some R programming techniques that one may find useful in reducing the computational cost.

Contents

Chapter 1

Introduction

1.1 Report Structure

Here in Chapter 1, the literature review in Section 1.3 will introduce necessary definitions of financial options and option pricing methods that are included in `rmcop` package. This is accompanied by some useful formula deductions that will be seen useful in Chapter 3, where we will discuss the programming implementation of these models.

In Chapter 2, we will introduce three existing option pricing packages in R (`derivmkts`, `fOptions`, `RQuantLib`), and explain their functionalities. We will also discuss their advantages and limitations.

In Chapter 3, we will discuss the package structure of `rmcop`. This includes: 1. a discussion on R's generic Object Oriented Programming (OOP) environment; 2. an outline of the package's structure; and 3. a detailed elaboration of the R implementation techniques of the financial models introduced in Section 1.3.

Finally, in Chapter 4, we will demonstrate examples and interpretations using `rmcop`. The report is then concluded by a discussion on `rmcop`'s limitations and plans for further development.

1.2 Package Description

The R package `rmcop` is developed for pricing financial options. The name "rmcop" stands for "R Monte Carlo Option Pricing". The package supports Monte Carlo-based pricing for European vanilla options and European exotic options, including Asian, Barrier, Binary, and Lookback options. It also provides some deterministic methods for both European and American vanilla options pricing, including option pricing via the Binomial lattice tree and Black-Scholes formula.

Unlike most existing packages in the R Ecosystem, `rmcop` is developed through an object-oriented approach. User's argument inputs for pricing functions are encapsulated in corresponding objects, and pricing functions are themselves methods. This introduces ordered arguments control and easier functions application.

1.3 Literature Review

1.3.1 Pricing of Financial Options

An option is a common financial derivative[1] in the market. Options can be roughly classified into two types, "call" or "put". A call option, "gives its holder the right (but not the obligation) to purchase from the issuer a prescribed asset[2] for a prescribed price (strike price) at a prescribed time (maturity/expiration date) in the future.[1]" Similarly, a put option gives the right to sell an asset at the strike price at maturity.

There are two most common styles of financial options [1]. A European option can only be exercised (i.e. executed) at maturity, whereas an American option can be exercised at any time prior to maturity (which is a more complex setting). Generally, an American option gives more freedom of exercising, so it is usually priced higher than a European option, giving all other parameters being equal.

Here is a hypothetical example. Suppose you hold a European call option on Apple's stock with a strike price of 150 and a maturity of 1 year. This option gives you the right to purchase a share of Apple's stock at exactly 150 on the day one year later. Let's assume that after a year, the market price of a share of Apple's stock is 160, and with the option in hand you, can purchase the stock with $160 - 150 = 10$ less. The result of £10 would be the benefit, or the **payoff**, that the option provides you. On the contrary, if after a year, the stock price is 140, which is less than the strike price of 150, there is no point exercising the option and purchasing the stock at a higher price. So the payoff of the option, in this case, would be zero.

The most typical "family" of options are the so-called "vanilla options" [1], which include no special or unusual features. For vanilla cases, a European call option has a payoff of $(S_T - K)^+ := \max(S_T - K, 0)$, and a European put option has a payoff of $(K - S_T)^+ := \max(K - S_T, 0)$ [2] at the point of maturity T in time. Here, S_t is the underlying asset price at some future time $t \in [0, T]$, and K is the option's strike price.

Options with any "special or unusual features" are called "exotic options". The exotic options whose pricing is supported by `rmcop` package is introduced below based on the definitions given in *An Introduction to Financial Option Valuation* [1].

- **Asian Options** Asian options' payoff are determined by the average price of the underlying asset throughout its life span, denoted $\bar{S}$.

  - An average price Asian call option has payoff at the expiry T given by $\max(\bar{S} - K)$.
  - An average price Asian put option has payoff at the expiry T given by $\max(K - \bar{S})$.
  - An average strike Asian call option has payoff at the expiry T given by $\max(S_T - \bar{S})$.
  - An average strike Asian put option has payoff at the expiry T given by $\max(\bar{S} - S_T)$.

- **Barrier Options** Barrier options have a payoff that switches on or off depending on whether the asset crosses a pre-defined level (barrier) B.

  - A down-and-out (knock-out) call option has a payoff that is zero if the asset crosses some predefined barrier $B < S_0$ at some time in $[0, T]$. If the barrier is not crossed then the payoff becomes that of a European call, $\max(S_T - K, 0)$.
  - A down-and-in (knock-in) call option has a payoff that is zero unless the asset crosses some predefined barrier $B < S_0$ at some time in $[0, T]$. If the barrier is crossed then the payoff becomes that of a European call, $\max(S(T) - K, 0)$.

- **Binary Options** A binary (a.k.a. cash-or-nothing) option have payoff at expiry being either some specified value A or zero.

---

[1]Security whose value depends on the value of an underlying (i.e. the related) asset.
[2]Underlying asset/stock, the word "asset" and "stock" may be used interchangeably through this report.

– A binary call option has payoff A if $S_T > K$ and zero otherwise.

– A binary put option has payoff A if $S_T < K$ and zero otherwise.

- **Lookback Options** The payoff for a lookback option depends upon either the maximum $S^{\max}$ or the minimum value $S^{\min}$ attained by the asset throughout the price trajectory.

    – A fixed strike lookback call option has payoff at expiry T given by $\max(S^{\max} - K, 0)$.

    – A fixed strike lookback put option has payoff at expiry T given by $\max(K - S^{\min}, 0)$.

    – A floating strike lookback call option has payoff at expiry T given by $S_T - S^{\min}$.

    – A floating strike lookback put option has payoff at expiry T given by $S^{\max} - S_T$.

1.3.2 Brownian Motion

A *stochastic process*, by definition [3], is a family of random variables $\{X_\gamma, \gamma \in \Gamma\}$ defined on $\Omega \times \Gamma$ taking values in $\mathbb{R}$. Thus, the random variables of the family (measurable for every $\gamma \in \Gamma$) are functions of the form:

$$X(\gamma, \omega) : \Gamma \times \Omega \mapsto \mathbb{R}$$

Where $(\omega, \mathcal{A}, P)$ is a probability space described by samples ω, action space $\mathcal{A}$, and probability measure P.

For $\Gamma = \mathbb{N}$, we have a *discrete process*, and for $\Gamma \subset \mathbb{R}$, we have a *continuous process*. In our context, we will consider Γ to represent the scope in time and takes values from 0 to maturity T.

A *stochastic process* $\{X(t), t \geq 0\}$ is said to be a *Brownian motion* if [3]:

1. $X(0) = 0$;

2. $\{X(t), t \geq 0\}$ has stationary and independent increments.

3. for every $t > 0$, $X(t) \sim N(0, \sigma^2 t)$.

In simple terms, a Brownian motion starts at a point in time, and for each time step, we move to a new point which is determined by a distribution centered on the current location.

A *Standard Brownian Motion* (i.e. Wiener Process, abbreviated as SBM) inherits the properties above, and has a volatility measure of zero (i.e. $\sigma = 0$). So for an SBM W, we have $W(t) \sim N(0, t)$.


Figure 1.1: Trajectories of Wiener Process

Figure 1.1 demonstrates the trajectories of a one-dimensional SBM, where each grey line represents a simulated path of a Wiener process.

Notice that the expectation of an SBM at any time $t \geq 0$ is zero, this makes an SBM a *martingale*, which is a stochastic process that "does not tend to rise or fall[3] [4]". If S_t is the value of an SBM at time t, assume $S_0 = 0$, its SBM dynamics can be defined by the Stochastic Differential Equation (SDE):

$$dS_t = S_t dW_t$$

We can generalise the scenarios by considering drifts through time (measured by μ) and scaling the amount of diffusion with respect to time (measured by σ). For S_t that follows a *Geometric Brownian Motion* (GBM), its dynamics can be described by the SDE [4]:

$$dS_t = \mu S_t dt + \sigma S_t dW_t$$

Where $\mu S_t dt$ is called the *drift term* and $\sigma S_t dW_t$ the *diffusion term*.

Expanding on S_t, for a function $f : \mathbb{R} \mapsto \mathbb{R}$, the GBM $f(S_t)$ is described as, according to Itô's rule [4]:

$$df(S_t) = \left[\mu_t \frac{\partial f}{\partial S}(S_t) + \frac{1}{2} \frac{\partial^2 f}{\partial S^2}(S_t) \right] dt + \sigma_t \frac{\partial f}{\partial S} dW_t$$

Using the Itô-Doeblin rule for $f : \mathbb{R} \mapsto \mathbb{R}$ [4], we can derive the following explicit solution to the SDE:

$$f(S_t) = f(S_0) + \int_0^t \mu_t \frac{\partial f}{\partial S}(S_t)dt + \frac{1}{2} \int_0^t \sigma_t^2 \frac{\partial^2 f}{\partial S^2}(S_t)dt + \int_0^t \sigma_t \frac{\partial f}{\partial S}(S_t)dW_t \qquad (1.1)$$

Equation 1.1 will be shown useful in the Black-Scholes model introduced in the following section.

1.3.3   Black-Scholes Model

The first quantitative method in option pricing (perhaps in all finance) was developed by a French mathematician, Louis Bachelier, in 1900. In his work *The Theory of Speculation* [5], he deducted deterministic formulas for pricing European (vanilla) calls and put options as follows:

$$C(S,T) = SN(\frac{S-X}{\sigma\sqrt{t}}) - XN(\frac{S-X}{\sigma\sqrt{t}}) + \sigma\sqrt{t}N(\frac{S-X}{\sigma\sqrt{t}})$$
$$P(S,T) = XN(\frac{S-X}{\sigma\sqrt{t}}) - SN(\frac{S-X}{\sigma\sqrt{t}}) + \sigma\sqrt{t}N(\frac{S-X}{\sigma\sqrt{t}})$$

Being the earliest approach, Bachelier's formula had already outlined the relationship between the option price and asset price S, strike price X, and volatility measure σ, which are essential fragments in modern formulas. However, based on limited data, Bachelier's solution was built under some unrealistic assumptions [5]. The normality assumption violates the non-negativity of the stock price, and the formula's discrete measure in time omitted the effect of continuous

---

[3]That is, the expectation $E[W(t+k)] = W(t)$ for any increment in time k.

movements in the stock price. Also, the formula did not discount the effect of interest rates. These errors cause Bachelier's model fails to price options accurately.

It wasn't until the 1960s had further improvements been made to quantitative option pricing. In 1961, Case Sprenkle [6] introduced the Sprenkle formula. The formula addressed the above issues by describing the stock price by the more suitable log-normal distribution[4] and discounting for the effect of interest rate, which successfully explained the time value of an option. In the following decade, improvements have been made by scholars such as Boness and Samuelson [7], who introduced empirical constants to increase the effectiveness of Sprenkle's model.

The model was finalised by Black and Scholes in 1973 [7], who explained the stock price movement by GBM. The GBM was described in the form of a Stochastic Differential Equation (SDE), which effectively models the continuity of price movement. The solution (derived through Itô's lemma) of the Black-Scholes formula under a risk-neutral approach[5] eliminates the empirical measure of Sprenkle et al's model, which resulted in an objective and deterministic estimation of the option price, as is used by most contemporary pricing methods.

Recall the explicit form of an SDE $f(S_t)$ given by the Itô-Doeblin rule in Equation 1.1. We assign $f(S_t) = \log(S_t)$[6] where log is the natural logarithm. The solution is:

$$\log(S_t) = \log(S_t) + (\mu - \frac{\sigma^2}{2})t + \sigma W_t \tag{1.2}$$

Because $W_t \sim N(0, t)$, we can see that Equation 1.2 shows that:

$$\log(S_t) \sim N(\log(S_0) + (\mu - \frac{\sigma^2}{2})t, \sigma^2 t)$$

This implies that S_t follows a *log-normal* distribution. A log-normal distribution takes only a positive value, which addressed the negativity issue caused by Gaussian stock price models, like the one developed by Bachelier [5].

By taking the exponential on the two sides of the equation, one can obtain the following expression:

$$S_t = S_0 \exp \left\{ (\mu - \frac{\sigma^2}{2})t + \sigma W_t \right\} \tag{1.3}$$

Which describes the stock price S_t at anytime $t \in [0, T]$. When simulating real market conditions, one may replace the factor μ with r and q, which stands for the (fixed) interest rate and dividend yield rate, respectively.

$$S_t = S_0 \exp \left\{ (r - q - \frac{\sigma^2}{2})t + \sigma W_t \right\}$$

Under the risk-neutral assumption, the payoff of a (European vanilla) call option with strike price K and expiration T is given by the expectation $E[e^{-rT}(S(T) - K)^+]$, and the corresponding payoff of a (European vanilla) put option at the same strike price and expiration is given by $E[e^{-rT}(K - S(T))^+]$. Using the solution of the Black-Scholes SDE we can deterministically evaluate the payoffs as follows [1]:

---

[4]An log-normal distribution is defined on $(0, \infty)$, which addressed the non-negativity issue.

[5]The risk-neutral approach, in simple terms, is constructing a portfolio at a moment in time such that the portfolio value will be identical at the next moment in time regardless of the price movement, so the portfolio will be riskless to the price movement.

[6]The transformation constructing the log-normal distribution.

$$C(S(0), T) = S(0)\Phi(d_1) - e^{-rT}K\Phi(d_2)$$
$$P(S(0), T) = e^{-rT}K\Phi(-d_2) - S(0)\Phi(-d_1)$$

Where Φ is the cumulative normal distribution, $d_1 := \frac{\log(S(0)/K) + (r + \frac{1}{2}\sigma^2)T}{\sigma\sqrt{T}}$, and $d_2 = d_1 - \sigma\sqrt{T}$. The Equations 1.4 and 1.5 are the so-called Black-Scholes formula for European options' price.

A generalised solution which addressed the impact of the presence of dividends (assuming the option of interest has an annual dividend yield rate of δ) specified as [2]:

$$C(S(0), T) = e^{-\delta t}S(0)\Phi(d_1) - e^{-rT}K\Phi(d_2) \tag{1.4}$$
$$P(S(0), T) = e^{-rT}K\Phi(-d_2) - e^{-\delta t}S(0)S(0)\Phi(-d_1) \tag{1.5}$$

Where now $d_1 := \frac{\log(S(0)/K) + (r - \delta + \frac{1}{2}\sigma^2)T}{\sigma\sqrt{T}}$, and d_2 has the same relationship described above.

A further result of the American option cases is such that an American call option is never optimal (i.e. the estimated value would be the same as the European call option under the same conditions); and that an American put option's value does not have an analytical form and required numerical methods to calculate[7].

1.3.4   Binomial Lattice Model

In 1979, based on the risk-neutral approach used by the Black-Scholes model, Cox, Ross and Rubinstein (CRR) [8] introduced a Binomial lattice tree model for modelling stock prices.

To construct a Binomial tree of stock prices, one breaks down the option's life $[0, T]$ into n time steps with fixed interval $\Delta t = T/n$. For each time step, the stock price can move either up by a factor u with probability $\hat{p}$ or down by a factor d with probability $\hat{q} = 1 - \hat{p}$.

Figure 1.2: Binomial Lattice Tree for Stock Price with 2 Steps

As Figure 1.2 shows, we will obtain a total of $n + 1$ nodes at maturity T. For each node at the final step, we can obtain the corresponding payoff of an option related to that stock price. For example, recall from the above section, the $n + 1$ possible payoffs $P(S(T))$ of an vanilla European call option would be:

---

[7]This is related to the studies on Monte Carlo pricing for American Options, which is beyond the current scope of `rmcop`.

$$P(S(T)) = (S(T) - K)^+ \tag{1.6}$$

The probability $\hat{p}$ is set to be "risk neutral" in a way such that:

$$
\begin{aligned}
S(t_i) &= e^{-r\Delta t}\mathbb{E}[S(t_{i+1})] \\
&= e^{-r\Delta t}[\hat{p}uS(t_i) + \hat{q}dS(t_i)], \quad \text{for } i = 0, ..., n-1
\end{aligned}
$$

And it can be calculated using the formula [8]:

$$\hat{p} = \frac{e^{r\Delta t} - d}{u - d} \tag{1.7}$$

This means the current node's stock price is equal to the expected value of the two future nodes it can go in the next time step, discounting the interest. By this risk-neutral setup, a portfolio which consists only of this stock would have a fixed return independent from price movement.

Under this setting, the option payoff at each node is computed in a similar way. Suppose payoff P_{ij} denotes the payoff of the jth node at time step i, it can be computed via recursive form:

$$P_{ij} = e^{-rT}[\hat{p}P_{(i+1),j+1} + \hat{q}P_{(i+1),j}] \tag{1.8}$$

As mentioned above, the payoff of an option at maturity T is obvious to obtain. By applying this recursive method forward from $t = T$ back to $t = 0$, we will obtain a deterministic estimate of the current payoff (fair value) of the option of interest.

1.3.5 Monte Carlo Option Pricing

As the financial market develops and more complicated options emerge, in many realistic cases, one cannot find a deterministic solution for pricing options. However, thanks to the advancement of computer power, one can simulate price trajectories for enormous times, and estimate the payoff of the option of interest by simply taking the average of the option payoff under each simulated trajectory. This method is known as the Monte Carlo method. The very first attempt at applying a computational method in option pricing is by Phelim Boyle in 1977. A more serious (and effective) approach was introduced by Paul Glasserman in the 1990s based on the Black-Scholes model. Until now, the field of Monte Carlo option pricing is still under active development and is widely used by "quants". The contents below will elaborate implementation method as the one introduced in Glasserman's *Monte Carlo Methods in Financial Engineering* [2].

We know that a GBM is a continuous process, and its dynamics are described by an SDE with respect to the t through infinitesimal dt. In order to simulate the process, we need to discretise dt to the computable Δt.

Recall the explicit form of the GBM (Equation 1.3) introduced in the Black-Scholes model, we can expand on it to describe the change in asset price $S_{t+\Delta t}$ as:

$$S_{t+\Delta t} = S_t \exp\left\{(\mu - \frac{\sigma^2}{2})\Delta t + \sigma W_{\Delta t}\right\}$$

Knowing that $W_t \sim N(0, t)$, suppose we have $Z :\sim N(0, 1)$, we can discretise $W_t = Z\sqrt{t}$. So we can modify the above equation to:

$$S_{t+\Delta t} = S_t \exp\left\{(\mu - \frac{\sigma^2}{2})\Delta t + \sigma Z \sqrt{\Delta t}\right\}$$

Suppose we discretise the continuous time interval $[0, T]$ into m time steps with fixed length $\Delta t = \frac{T}{m}$. We can derive the iterative formula:

$$S_i = S_{i-1} \exp\left\{(\mu - \frac{\sigma^2}{2})\Delta t + \sigma Z_i \sqrt{\Delta t}\right\}, \qquad i = 1, ..., m \tag{1.9}$$

We can also modify the equation to address the effect of interest rate and dividend yield rate according to what's mentioned below Equation 1.3:

$$S_i = S_{i-1} \exp\left\{(r - q - \frac{\sigma^2}{2})\Delta t + \sigma Z_i \sqrt{\Delta t}\right\}, \qquad i = 1, ..., m \tag{1.10}$$

By simulating m standard normal random variables $Z_1, ..., Z_m$, we will be able to generate a full trajectory of the price movements. We will represent this using the below pseudocode [2].

for $i \leftarrow 1, ..., m$ **do**
$\quad Z_i \leftarrow N(0, 1)$
$\quad S_i \leftarrow S_{i-1} \exp\left\{(\mu - \frac{\sigma^2}{2})\Delta t + \sigma Z_i \sqrt{\Delta t}\right\}$
end for

Noticing that as $W_T \sim N(0, T)$, and that individual increments for Brownian motions are independent, the approximation of W_T using Z_i's is lossless (i.e. the number of time steps/length of Δt, has no effect on how well the approximation will be). For path-independent options[8], it may be wise to take $m = 1$, i.e. $\Delta t = T$ to minimise the computational cost.

If we are to simulate n times, for the ith trajectory, $i = 1, ..., n$, one may employ formulas introduced in List 1.3.1 and discounting the effect of interest rate and dividend rate to obtain the corresponding discounted payoff (i.e. fair price) of the option, C_i. For example, the discounted payoff of a Vanilla call option, at maturity, can be calculated as:

$$C_i = e^{-rT}(S_T - K)^+ \tag{1.11}$$

The Monte Carlo estimator for the option price is then simply the arithmetic mean $\hat{C}_n := \frac{1}{n}\sum_{i=1}^{n} C_i$. This estimator is *unbiased* and *strongly consistent* [2].

For large enough n, we can construct a confidence interval using the standard error, given by [2]:

$$SE_C = \frac{s_C}{\sqrt{n}} \tag{1.12}$$

$$s_C = \sqrt{\frac{1}{n-1}\sum_{i=1}^{n}(C_i - \hat{C}_n)^2} \tag{1.13}$$

Glasserman in his work [2] also indicated the criterion for the simulation estimators' efficiency: computing time, bias, and variance. The estimators considered in our package are unbiased, and variance reduction techniques are beyond our scope. So we will only discuss some R programming methods in reducing the computing time to increase estimation efficiency, as we will see in Chapter 3.

---

[8]Options whose payoff only depends on the stock price at maturity, S_T.

Chapter 2

Existing Option Pricing Packages within R Ecosystem

2.1 Introduction

In this chapter, we will introduce three R packages on CRAN which are used for pricing financial options. We will comment on each of their strengths and weaknesses.

2.2 Packages Review

2.2.1 derivmkts

The R package `derivmkts` [9], which abbreviates for "derivative markets", provides a collection of functions for pricing and analysing financial derivatives. It includes options pricing methods, including Monte Carlo simulations and binomial tree method, as well as other market analysis tools including calculating Greeks, implied volatility, and so on.

The package `derivmkts`'s functions are functional-specific and relatively easy to use. For instance, its `asianmc` functions provide pricing calculations exclusively for Asian options using the Monte Carlo method, and `perpetual` functions are specialised for pricing perpetual options[1] only. Implementing these functions is simple, but this high function independence lacks expandability. Without reusing existing functions, new functions have to be created bottom-up. As a result, the package only provides limited functions.

On the other hand, the package document is sparse and sometimes incomplete. Algorithms such as the `binom` for generating the Binomial tree and `simprice` for running Monte Carlo price simulations have not been benchmarked against other sources and are computationally inefficient. For example, the below codes benchmarked price trajectories sampling using `simprice` from `derivmkts` against the function in `rmcop` using the package `microbenchmark` [10].

```
1   func.rmcop <- function() { # Monte Carlo pricie simulation using rmcop
2       op <- option("european", type = "call", K = 40, t = 0.25)
3       op.env <- option.env(S = 40, r = 0.08, q = 0, sigma = 0.3, n = 100,
            steps = 100)
4       price.option(op, op.env)
5   }
6
7   func.derivmkts <- function() { # Monte Carlo pricie simulation using
        derivmkts
```

---

[1]Options with maturity $T = \infty$, i.e. an option that never matures.

```
 8        simprice(s0 = 40, v = 0.3, r = 0.08, tt = 0.25, d = 0,  trials =
          100, periods = 100, jump = FALSE)
 9    }
10
11    microbenchmark::microbenchmark(
12        func.rmcop(),
13        func.derivmkts()
14    )
```

The results are:

```
 1    > Unit: microseconds
 2    >             expr      min       lq      mean   median        uq       max
      neval cld
 3    >      func.rmcop()    859.9    951.0  1040.669   985.4   1055.45    2985.8
      100  a
 4    > func.derivmkts() 28316.3 28691.1 31946.188 29055.6 29936.10   74694.6
      100  b
```

From which we can see that the average computing time of simulating 100 trajectories with 100
times steps on each trajectory is over 30 times longer for `derivmkts` than `rmcop`.

The reason is because `derivmkts`, as a supplement programme for educational purposes, had its
functions coded in a very direct way corresponding to the underlying formulae. This results in
a pervasive use of for-loops in R without any specific adjustments to improve the computation
speed.

2.2.2 fOptions

Similar to `derivmkts`, `fOptions` [11] is also a supplement package for educational purpose. It
provides a collection of functions for pricing and analyzing various financial options, such as
European, American, Asian, exotic, and barrier options. The package also includes methods
for calculating implied volatility, greeks, and binomial trees. It is useful for researchers and
practitioners who want to apply option pricing models to real-world data and scenarios.

In contrast to `derivmkts`, it has a clearer documentation and examples. However, the package
functions still lacks computing efficiency (having the same issue as the `derivmkts`). A more crucial
issue with the package is that it is no longer supported by the current R environment, and its
development was terminated since 2020.

2.2.3 RQuantLib

Both `derivmkts` and `fOptions` are supplementary materials for educational purpose. `RQuantLib`
[12], on the other hand, is a professional tool for industrial computational finance.

`RQuantLib` is an R interface to the QuantLib software, an open-source library for quantitative
finance based on C++. The purpose of RQuantLib is to enable R users to perform complex
financial calculations and simulations using the functionality of QuantLib. Some of the strengths
of RQuantLib are:

- It covers a wide range of financial instruments and models, such as bonds, options, swaps, interest
rate models, volatility models, etc. - It allows for easy integration with other R packages and data
sources, such as time series analysis, graphics, optimization, etc. - It offers high performance and
reliability due to the use of C++ code and rigorous testing.

However, due to its highly professional nature, it requires pre-installation and configuration of QuantLib software on the system, which may be challenging for many users. Also, because it is a supplement of the QuantLib library, it's own documentation and examples are very limited compared to other R packages, making it not a user-friendly choice for less-intensive option pricing purpose.

Chapter 3

Package Development

3.1 Introduction

In this chapter, we will elaborate on the development of `rmcop`. We will first present the general structure of the package, briefing on the purposes of package files and functions. Then, we will discuss the R implementation of pricing models mentioned in the literature review in Section 1.3 in detail.

3.2 Machine Environment

The majority of the development and testing of the package is completed using my laptop. The configuration is briefed below.

CPU: AMD Ryzen 7 4800H with Radeon Graphics (base 2.90GHz)
RAM: 16(15.4)GB DDR4 3200MHz
Max TDP: 45W

The setup is typical for personal computing devices, so the benchmarking and profiling presented in the report are reasonable references for ordinary package users.

We have used the most up-to-date software setup until March 2023. The running environment is R 4.2.2 on RStudio 2022.12.0. The package framework is constructed using package `devtools` [13] version 2.4.5.

Until the time drafting this report, exported functions have passed the test using the R package `testthat` [14] version 3.1.7. The testing functions is structured in the following framework.

```
1  test_that("test title", {
2      expect_equal(<operation>,
3                   <expected result>,
4                   tolerance = <precision>)
5  })
```

Using the `expect_equal` setup, the test will only pass if the operation's output is equal to the expected result under the specified level of precision.

For the package's deterministic methods, the binomial tree and the Black-Scholes formula pricing, the tolerance is set to `1e-5` (i.e. 1×10^{-5}) and `1e-7`, respectively. For Monte Carlo approximation, tolerance is set to a higher `1e-1` to check convergence.

3.3 Package Structure

The overlying workflow of using `rmcop` package to price financial option is displayed in Figure 3.1 below. One shall see that with the OOP framework, we encapsulate most arguments into two classes of objects: `option` and `env`.
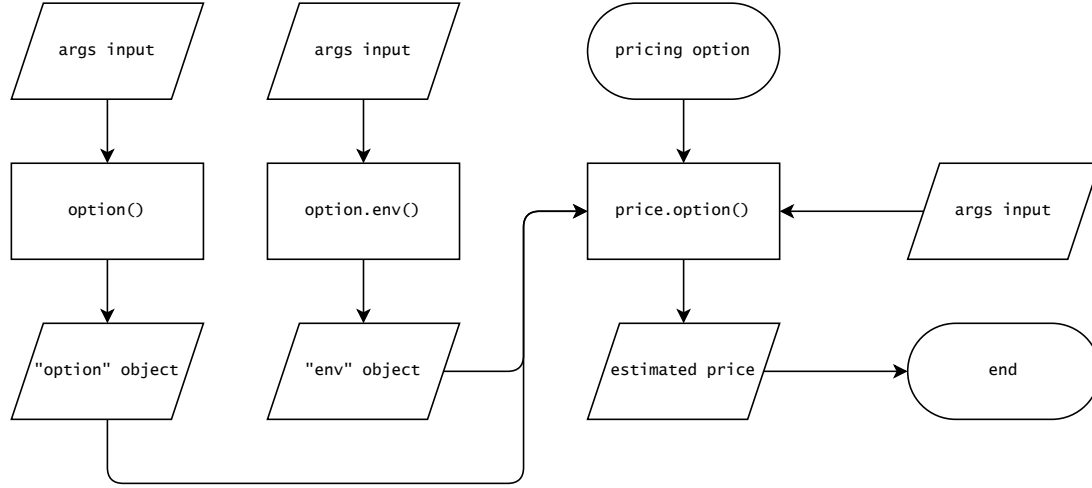

Figure 3.1: Flowchart of option pricing using rmcop package

As Figure 3.1 indicates, there are only three exported[1] functions. Their purpose is shown in Table 3.1 below:

Function	Description
`option()`	Create new `"option"` class object, which encapsulates the characteristics of an option
`option.env()`	Create new `"env"` class object, which encapsulates the market environment setup
`price.option()`	Pricing the option based on specified option, market environment, and method input

Table 3.1: Functions to Call

To keep the functions (exported and internal) organised and to make each script restricted to a maintainable length, they are divided into 7 R scripts listed below:

File	Contents
`Option.R`	Methods creating and updating option and option.env objects
`Price.R`	Pricing functions takes objects input and calls specific pricing engines
`MonteCarlo.R`	Monte Carlo option pricing method engine
`BlackScholes.R`	Black-Scholes option pricing method engine
`Binomial.R`	Binomial option pricing method engine
`Tools.R`	Supplementary functions used in the package

Table 3.2: Package R Scripts

`Option.R` and `Price.R` contains all constructor functions for S3 class objects. `Price.R` contains the overlying pricing functions which will trigger the pricing engine functions that are contained

---

[1] When installing an R package, only exported functions can be directly called and executed by the user; other "unexported" functions might serve as internal components of the package.

in `MonteCarlo.R`, `BlackScholes.R`, and `Binomial.R`. `Tools.R` includes supplementary functions, including input integrity checking and Monte Carlo price trajectories plotting.

3.4 Object Oriented Programming

3.4.1 R Framework for OOP

Existing packages in the R ecosystem provide comprehensive pricing algorithms for financial derivatives. However, their functions are used in a Procedural Oriented (POP) way. POP functions can be called directly by passing in required arguments. For simple options pricing cases, such as pricing individual options, using POP functions is intuitive. However, many real scenarios require pricing options in a complex formulation, such as compounded options (i.e. options with underlying assets being another option) and a combination of options (i.e. spread, straddle, stringle, and other option strategies). In these situations, managing numerous arguments for POP functions can be difficult and inefficient.

The `rmcop` package proposed an Object Oriented Programming (OOP) approach for pricing financial options in R. It encapsulates multiple arguments, such as option style, type, strike price, and maturity time, into an `option` class object. It also allows encapsulation of other market environments arguments, such as interest rate, dividend yield rate, and volatility measure into an `option.env` class object. The OOP structure enables easier variables managements and facilitates the comparison of prices among different sets of options and market environments.

The codes below demonstrate the calculation of a theoretical European vanilla call option with strike price $K = 20$, maturity $t = 0.75$, under the market condition such that the current price is $S = 20$, the fixed interest rate is $r = 1\%$, and market volatility measured by $\sigma = 0.1$. We use Monte Carlo (i.e. `mc`) method with $n = 100$ replications and number of time steps per replication $steps = 1$.

The OOP approach requires extra steps declaring objects before using the function, but once the declaration is completed, calling the pricing function is much simpler than the POP approach. As above, it only required two (object) arguments, `obj` and `env`.

R provides two major ways to perform OOP, the S3 and S4 methods. The S3 class objects are based on R `list` objects. An R list contains a `class` attribute, which can be customised into a string (or vector of strings) that can be interpreted as a list's corresponding S3 class (i.e. so that the list itself became an object of that class). Other items within the list can be treated as object properties under the OOP scope. Here is an example of how to implement OOP using the S3 method in R.

```r
student <- function(age, gender) {
    obj <- list(
        "age" = age,
        "gender" = gender
    )
    class(obj) <- "student"
    obj
}

John <- student(
    "age" = 20,
    "gender" = "male"
)

John
> $age
```

```
17   > [1] 20
18   >
19   > $gender
20   > [1] "male"
21   >
22   > attr(,"class")
23   > [1] "student"
```

We first define a new `list` object named "John", which contains two items/properties, age and gender. Then, we redefined the class of this list using the `class()` function to "student". Now, we have created a new object named "John" of the class "student" under the S3 scheme.

R's S3 method also provides a naive way of expressing inheritance, which is by expressing the `class` attribute of an object with a vector instead of a single string. Continuing our previous example, let's now create a new object "Sara" using the following code:

```
1    international.student <- function(age, gender, nationality, language) {
2
3        obj <- student(age, gender)
4        obj <- c(
5            obj,
6            "nationality" = nationality,
7            "language" = language
8        )
9        class(obj) <- c("international", "student")
10       obj
11   }
12
13   Sara <- international.student(
14       age = 21,
15       gender = "female",
16       nationality = "Singapore",
17       language = "Chinese"
18   )
19
20   Sara
21   > $age
22   > [1] 21
23   >
24   > $gender
25   > [1] "female"
26   >
27   > $nationality
28   > [1] "Singapore"
29   >
30   > $language
31   > [1] "Chinese"
32   >
33   > attr(,"class")
34   > [1] "international" "student"
```

Here, we see that the object representing Sara not only contains the properties a "student" class object has but also contains two extra properties "nationality" and "language". In this case, the objects from class `international student` are inherited from the father class `student`, which is

reflected by the `class` attribute being a vector and has the last item named "student". To check whether an object is from a class, we can use the `inherits` function. For example:

```
1  inherits(Sara, "student") # True
2  inherits(John, "student") # True
3  inherits(Sara, c("international", "student"), which = TRUE) # 1 1
4  inherits(John, c("international", "student"), which = TRUE) # 0 1
```

The S4 method requires more rigorous class and object definition, as one would typically see in an OOP language such as Python and Java. The development of our pricing functions does not require a rigorous OOP structure or defining generic functions, so the S3 framework is sufficient for its development.

3.4.2   Option and Environment Class Objects

As we mentioned in Section 3.3, `rmcop` requires users to encapsulate most of their input arguments as two classes of objects, `option` and `env`. Thus, we first need to build functions (constructors) that allow users to create/initialise new objects.

An option class object should contain basic properties that define a financial option. This includes the option's style, type, strike price K, and maturity T. These properties are common for both vanilla and exotic options, so we should encapsulate them by the overlying class.

```
1   option <- function(style, type, K, t) {
2       obj <- list(
3           "style" = style,
4           "type" = type,
5           "K" = K,
6           "t" = t
7       )
8       class(obj) <- "option" # Specifying class attribute to the S3 object
9       obj
10  }
```

Exotic options usually requires additional properties. For instance, a binary option require an argument which specifies the fixed payoff it generates; similarly, a barrier option require two additional properties to specify a price barrier and to classify whether it is a "knock-in" or "knock-out" option, respectively.

To address these differences, creating "sub-classes" and using inheritance techniques via OOP framework is an intuitive and elegant approach. We shall update our previous function by including a new argument named `option`, which allows user to specify the family/name of the (exotic) option the user wish to define[2].

```
1   option <- function(style, option = "vanilla", type, K, t, ...) {
2
3       obj <- list(
4           "style" = style,
5           "type" = type,
6           "K" = K,
7           "t" = t
8       )
9
```

---

[2]The default value is set to `"vanilla"`, since this is the most common family of options.

19

```
10        # Direct to corresponding sub-class constructors
11        if (option == "vanilla") {
12            obj <- option.vanilla(obj)
13        } else {
14            obj <- option.exotic(obj, option, ...)
15        }
16        obj
17    }
```

As shown by the codes above, depending on the value of input for the argument `option`, the constructor calls one of the two sub-class constructors. The function `option.vanilla` will take the list created with four basic properties and assign it with a `"vanilla" "option"`[3] class attribute. On the other hand, the function `option.exotic` will take the input for argument `option` together with additional user inputs, and return an object of class `"XXX" "option"`, where `XXX` is the family of option that the user specified. An example is demonstrated below:

```
1    option(style = "european",
2           option = "barrier",
3           type = "call",
4           K = 20,
5           t = 1,
6           barrier = 21,
7           is.knockout = TRUE)
8    > $style
9    > [1] "european"
10   >
11   > $type
12   > [1] "call"
13   >
14   > $K
15   > [1] 20
16   >
17   > $t
18   > [1] 1
19   >
20   > $barrier
21   > [1] 21
22   >
23   > $is.knockout
24   > [1] TRUE
25   >
26   > attr(,"class")
27   > [1] "barrier" "option"
```

In this example, the user specified the option to be a "barrier" option, which required additional parameters `barrier` and `is.knockout` (detailed discussion see sections below on individual pricing functions). As we can see, the `class` attribute became `"barrier" "option"`, which follows what we have discussed earlier.

After completing the constructor for the option class, we need to build the constructor for the environment class. For an object of the class `env`, we wish to encapsulate some market conditions that will be considered during our pricing procedure. The parameters we considered here are pricing method, current underlying asset price S, interest rate r, dividend yield rate q, volatility

---

[3]This is the printed form of the vector `c("vanilla", "option")`

measure *sigma*, and two supporting parameters `n` and `steps`. Following the standard procedure of creating an S3 object, we have the following algorithm:

```r
option.env <- function(method = "mc", S, r, q = 0, sigma, n = NULL,
steps = NULL) {
    env <- list(
        "method" = method,
        "S" = S,
        "r" = r,
        "q" = q,
        "sigma" = sigma,
        "n" = n,
        "steps" = steps
    )
    class(env) <- "env"
    env
}
```

Since the Monte Carlo method is the most commonly used pricing method (and the core of this package), we set the default value of the argument `method` to the abbreviation `"mc"`. The parameters of the current price, interest rate, and volatility measure are mandatory, as they are essential inputs for the pricing functions. The dividend yield rate is set to a default zero to cover the majority of non-dividend-bearing options. Finally, the two supplementary parameters `n` and `steps` are allowed to be left as `NULL` depending on whether the specified pricing method requires these arguments to be specified.

3.5   price.option

As we mentioned in Section 3.3, `price.option` is the overlying function in the script `Price.R`. It acts as a medium that connects user input and pricing algorithms, making the package structure more operable.

To do this, we require this function to: 1. take objects/arguments input from users, 2. trigger corresponding pricing engines based on specified method, and 3. return the estimation of the option price. We can do so starting from the codes below:

```r
price.option <- function(obj, env, method = env$method, n = env$n, ... ,
all = FALSE) {
    check.method(method)
    res <- get(paste0("price.option.", method))(obj = obj, env = env, n
    = n, ..., all = all)
    res
}
```

As we see, the input argument of the function are:

- `obj` the pre-defined option object.

- `env` the pre-defined environment object.

- `method` a string to determine which pricing method to use.

- `...` additional arguments that are required/provided by specific pricing methods.

- **all** a boolean[4] value to determine whether only a number of estimate should be returned or should all relevant data used in the computation should be returned as well.

Going through the code with more details, in line 2, `check.method` is a function in `Tools.R` which checks if a pricing method is specified and if it is supported by the package.

The following `get(<function name>)(<arguments>)` structure in line 3 executes the function with the name provided and takes specified argument inputs. This line triggers corresponding sub-functions that evoke specific pricing engines. For example, the sub-function for Binomial tree option pricing is shown below:

```
1   price.option.binomial <- function(obj, env, n, u, d, all) {
2       res <- get(paste0(class(obj)[1], ".binomial"))(obj = obj, env = env,
        n = n, u = u, d = d, all = all)
3       res
4   }
```

These sub-functions are all named in the format `"price.option.XXX"`, where `"XXX"` is the string defined by argument `method`. These sub-functions also used a similar structure to trigger corresponding pricing engines for a different family of options. The pricing engines are all named in the format `"family.method"`. For instance, the Monte Carlo pricing engine for the lookback option is named `lookback.mc`. Using the trick of `get()()` allowed us to decide which subordinate function to call without using lengthy if-else statements.

The input objects and method-specific arguments will then be passed to the pricing engine, where the information encapsulated within the objects will be extracted and used to compute the estimation of the option's price. The detailed construction of these engines will be discussed in the sections below.

3.6 Deterministic Methods

Deterministic methods provide closed-form formulas which build upon their relevant assumptions to price options. When pricing an option using a deterministic method, given the same initial conditions, the result estimates should always be the same.

The deterministic methods that are included in `rmcop` are the Black-Scholes formula and Binomial Lattice Tree. The package's implementation of them is limited to pricing American and European Vanilla options.

3.6.1 Black-Scholes Model

`vanilla.bs` is the pricing engine function for Black-Scholes formula. The solutions of call and put options are specified by Equation 1.4 and 1.5, respectively.

The implementation is simple, as we shall only "translate" the formula to R code. This can be done with the algorithm below:

```
1   d1 <- bs.d1(S, K, t, r, q, sigma)
2   d2 <- d1 - sigma * sqrt(t)
3
4   if (type == "call") {
5       price <- exp(-q * t) * S * pnorm(d1) - exp(-r * t) * K * pnorm(d2)
6   } else if (type == "put") {
```

---

[4]i.e. logical value, which evaluated to either true or false

```
7        price <- exp(-r * t) * K * pnorm(-d2) - exp(-q * t) * S * pnorm(-d1)
8    }
```

Where d_1 is computed by the function:

```
1    bs.d1 <- function(S, K, t, r, q, sigma) {
2        (log(S / K) + (r - q + sigma^2 / 2) * t) / (sigma * sqrt(t))
3    }
```

Which is the exact same formula that is specified under Equation 1.5. The engine will then simply return the variable `price` as the output.

Notice that by definition seen in Section 1.3, only a non-dividend bearing American call option can be priced using the Black-Scholes formula as well. This is because such options aren't worthwhile to be exercised earlier than maturity, so the resulting price will be the same as a European call with other arguments being identical. In other cases, an American option's price cannot be expressed in a closed form, and should thus be calculated using numerical methods.

To identify this issue and prevent users from pricing American-style options wrongly, we add the following code to generate a warning message:

```
1    if (style == "american") {
2        if (q != 0 | type != "call") {
3            warning("Black-Scholes formula not suitable, numerical
                 evaluation should be used instead.")
4        }
5    }
```

3.6.2   Binomial Lattice Tree

`vanilla.binomial` prices vanilla options via CRR's Binomial Lattice Tree method, as introduced in Section 1.3.

The R algorithm can be roughly divided into four segments: 1. compute the risk-neutral probability 2. generate binomial tree; 3. compute the payoff at the end step of the binomial tree; and 4. using backward recursive methods to compute the expected payoff of the option at the current time under the risk-neutral assumption.

First, after extracting the information from input objects, we can compute the risk-neutral probability using Euqation 1.7. This can be translated to the following code:

```
1    p <- (exp(r * dt) - d) / (u - d) # Compute riskless probability for an
         upward price movement
2    q <- 1 - p # Compute riskless probability for a downward price movement
```

Where variables `p` and `q` represents $\hat{p}$ and $\hat{q}$, respectively.

Second, to generate a binomial tree, we require four known parameters: the current underlying asset price S (or S_0), the ratio of an upward jump u, the ratio of a downward jump d, and the number of steps to compute n. By observing the Binomial tree through Graph 1.2, we see we can store all nodes of the tree into an $(n+1) \times (n+1)$ triangular matrix, as displayed in the following form:

$$S$$
$$dS \qquad uS$$
$$d^2S \qquad duS \qquad u^2S$$
$$\cdots$$
$$d^nS \qquad d^{n-1}uS \qquad d^{n-2}u^2S \qquad \cdots \qquad u^nS$$

Figure 3.2: Matrix Representing the Binomial Tree

As shown by Graph 3.2, each row of the lower triangular matrix represents a step of the binomial tree. To create such matrix shown in Figure 3.2, we can use the algorithm below:

```r
Binomial.tree <- function(S0, u, d, n) {
    dim <- n + 1
    S <- matrix(nrow = dim, ncol = dim) # Predefine the size
    temp <- c(S0)
    for (i in 1:dim) {
        S[i, ] <- c(temp, rep(NA, dim - i))
        temp <- c(temp[1] * d, temp * u)
    }
    S
}
```

We first build an empty matrix of the required size $(n+1) \times (n+1)$[5] and assign the current price S_0 to be the only element of row 1.

As we see from Graph 3.2, each row can be created via the upper row by: 1. Multiply the upper row's first element by the downward jump ratio d; 2. Multiply all elements (including the first) by an upward jump, and; 3. store these new values into a vector which is 1 unit longer than the previous row.

We encoded the above procedure using the for-loop shown in code lines 5-8 above. Here, `temp` is a record of each "upper row's" non-NA elements in the above context, and `c(temp, rep(NA, dim - i))` allows us to fill each new row's length to $n + 1$, so that we can directly update the rows in the predefined empty matrix.

A typical output of the `Binomial.tree` function is shown below:

```r
Binomial.tree(S0 = 8, u = 2, d = 0.5, n = 4)
>      [,1] [,2] [,3] [,4] [,5]
> [1,]  8.0   NA   NA   NA   NA
> [2,]  4.0   16   NA   NA   NA
> [3,]  2.0    8   32   NA   NA
> [4,]  1.0    4   16   64   NA
> [5,]  0.5    2    8   32  128
```

The third step is calculating the payoff at the terminal state of the binomial tree. In our setup, this would be the payoff calculated using the prices shown in the last row of the matrix. Recall the payoff of different types of options from Section 1.3, we will use an if-else statement to decide which payoff function to use:

---

[5]In R, this is represented by a matrix containing only `NA`s. Defining the size of the matrix ahead allows us to speed up the computation.

```
1   if (type == "call") {
2       fn <- sapply(Sn, function(x) {max(x - K, 0)}) # Payoff for call
        option
3   } else if (type == "put") {
4       fn <- sapply(Sn, function(x) {max(K - x, 0)}) # Payoff for put
        option
5   }
```

Here, variable `Sn` is a vector storing the last row of the constructed price matrix, and `fn` would represent the corresponding final step's payoffs. This will be used as an initial vector of our next recursive price valuation.

The fourth and final step is to use risk-neutral valuation to compute the current payoff from the final step's payoffs. For each step backwards in time, the recursive formula is shown by Equation 1.8. A simple pseudocode of the function encoding this formula would be:

> **procedure** RECURSE(step i payoffs)
>   **if** length of step i payoffs $= 1$ **then**
>   **return** payoff
>   **else**
>   Compute step i - 1 payoffs
>   **return** RECURSE(step i - 1 payoffs)
>   **end if**
> **end procedure**

For each layer of recursion, we will take the current step's payoff vector as input. Using Equation 1.8, the output of each recursion will be a vector that is 1 unit shorter than the input vector (meaning we are moving forward by one step in the binomial tree). When the length of the vector is finally reduced to 1, we will be reaching step 0, or the current time under the binomial tree setup. The result singular payoff will be the desired current payoff or the fair price of the target option.

The detailed implementation of the recursive method can be found in the package's source code under Appendix A.1.

Noticing that when the computer runs a recursive algorithm, the data will not be returned until the recursion reaches an end. This means the memory usage will accumulate continuously as the recurring layers stack deeper. Therefore, a binomial tree with steps $n = 1000$ will consume over 6GB of system memory and noticeably long time to complete, and a binomial tree with steps $n = 10000$ cannot be executed in R as the number of recurring layers will exceed the default limit (and RAM capacity of many personal devices).

To demonstrate this effect better, an example code is presented under Section 4.1.3.

3.7 Monte Carlo Methods

3.7.1 Simulating Price Trajectories

The core of Monte Carlo option pricing is simulating the price movements of the underlying asset. Recall the computable form of the Black-Scholes model (Equation 1.9) and the given pseudocode, we can directly implement this in R as follows.

```
1   # Setup arguments for calculation
2   t <- 2 # Expiration
3   n <- 100 # Number of trajectories to simulate
4   m <- 100 # Number of time steps per trajectory
```

```
5   S0 <- 20 # The current asset price
6   mu <- 0.05 # The drift coefficient
7   sigma <- 0.03 # The diffusion coefficient
8   dt <- t / m # The length of each time step
```

We will first directly implement the pseudocode in R, following Glasserman's framework [2].

```
1    mc.for <- function() {
2        S.mat <- matrix(nrow = m + 1, ncol = n)
3        for (j in 1:n) {
4            S <<- vector(length = m)
5            S[1] <<- S0
6            for (i in 2:(m + 1)) {
7                Z <<- rnorm(1)
8                S[i] <<- S[i - 1] * exp((mu + sigma^2 / 2) * dt + sigma * Z
                     * sqrt(dt))
9            }
10           S.mat[, j] <<- S
11       }
12   }
```

The resulting `S.mat` is a $(m+1) \times n$ matrix, where each column records a simulated price trajectory from $t = 0$ to $t = T$ of $m + 1$ steps (including the current stock price at time $t = 0$). We can plot the result trajectories using the `matplot()` function from the `graphics` package.

```
1    matplot(S.mat, type = "l", col = rgb(0, 0, 0, 0.2),
2            xlab = "dt", ylab = "price")
```


Figure 3.3: Monte Carlo Price Trajectories

From the result shown in Figure 3.3, we see a collectively upward movement for prices. This corresponds to our setup of a positive drift coefficient of 0.05.

However, the above code is not computationally efficient. R is relatively slow in running for-loops, and we can speed up the process by using the matrix method.

Recall Equation 1.2, where we have the explicit form of $\log(S_t)$. Using a similar discretisation

26

method, we can use the standard normal random variable Z to substitute W_t to form a recursive equation, such that:

$$\log(S_i) = \log(S_{i-1}) + (\mu - \frac{\sigma^2}{2})\Delta t + \sigma Z_i \sqrt{\Delta t}$$

Notice that assume Δt is fixed, the part $\log(\Delta S)_i := (\mu - \frac{\sigma^2}{2})\Delta t + \sigma Z_i \sqrt{\Delta t}$ is time independent. Therefore:

$$\log(S_{i+k}) = \log(S_{i-1}) + \sum_{j=i}^{k}\left[(\mu - \frac{\sigma^2}{2})\Delta t + \sigma Z_i \sqrt{\Delta t}\right] \tag{3.1}$$

$$= \log(S_{i-1}) + \sum_{j=i}^{k}\log(\Delta S_i) \tag{3.2}$$

For $i, k \in \mathbb{N}$ and $i + k \leq m$.

This allows us to reduce the two layers of for-loops into matrix operations.

```
mc.mat <- function() {
    Z <<- matrix(rnorm(n * m), ncol = n)
    logdS <<- (mu - sigma^2 / 2) * dt + sigma * Z * sqrt(dt)
    logS <<- log(S0) + apply(logdS, 2, cumsum)
    S <<- exp(logS)
    S <<- rbind(S0, S) # Add back the missing row of starting point S0
}
```

Here, we first define `Z` to be a matrix that contains all standard normal random variables required for a total of $n \times m$ number of steps. Then we define `logdS` as the matrix of all increments. Based on Equation 3.2, by using $\log(S_0)$ as the starting point and using cumulative sum `cumsum` function to add increment $\log(\Delta S)$ along steps of each trajectory, we will obtain the movement of prices in time. By some further post process shown above, we will get the same result as using the for-loops.

To test the difference in computing speed between the two methods, we will test the average running time using the R package `microbenchmark` [10]. The device configuration on which the test runs can be found in Section 3.2.

The `microbenchmark` function runs individual functions 100 times and records the summary statistics of the running times, as shown below:

```
microbenchmark::microbenchmark(
    mc.for(),
    mc.mat()
)

>     Unit: microseconds
>      expr      min       lq      mean    median       uq      max neval
> mc.for() 18327.2 19872.8 26257.504 25396.45 27853.6 60920.1    100
> mc.mat()   831.8    964.9  1200.425  1039.30  1094.5   8045.3    100
```

From the result, we see that `mc.mat` is substantially faster than `mc.for`. We should thus prefer the former to obtain better estimation efficiency.

The matrix method can also be implemented via the form shown in Equation 1.9 using cumulative product `cumprod` function. However, the running time for single multiplication is slightly longer than for addition. One may decide to use either method considering readability and efficiency.

The full look of the `mc.engine` is shown below:

```r
mc.engine <- function(type, K, t, S0, r, q, sigma, n, steps) {

    dt <- t / steps

    # Generate n random samples from N(0,1)
    Z <- matrix(rnorm(n * steps), ncol = n)

    # Get logarithm of price change per step
    increment <- (r - q - sigma^2 / 2) * dt + sigma * sqrt(dt) * Z

    # Use vectorised MC method to compute logarithm of S(t) at each time
    step
    logS <- log(S0) + apply(increment, 2, cumsum)
    S <- exp(logS) # Obtain the simulated stock price by taking
    exponential
    S <- S * exp(-q * t) # Adjust the asset price for dividends
    S <- rbind(S0, S) # Add column of current price
    S
}
```

Notice that the argument specifying the number of time steps to simulate per trajectory, m, is now specified under the name `steps`.

Here, the drift coefficient is decomposed to (fixed) annual interest rate r and (fixed) annual dividend yield rate q according to Equation 1.10. The number of steps in time m is directly named as `steps`.

Adapting this formulation, `rmcop` has a function named `mc.engine`, which takes the mentioned input arguments and returns a $(m+1) \times n$ matrix recording all generated trajectories. The rest of the functions described in the following subsections take `mc.engine` as an internal component and use the returned matrix to perform further calculations of option payoff.

Using R package `profvis` [15], we can profile the run-time of the function `price.option` under a reasonably time-consuming Monte Carlo setup. Here we will simulate 10000 trajectories, each with 1000 time steps. The result trajectory matrix will contain $10000 \times 1000 = 10000000$ simulated prices. The code for benchmarking is shown below, and the output is presented in Figure 3.4.

```r
# Creating option and environment objects
op <- option("european", "vanilla", "call", 20, 0.75)
op.env <- option.env(S = 20, r = 0.03, q = 0.01, sigma = 0.05)

# Profiling
profvis::profvis(
    price.option(op, op.env, n = 10000, steps = 1000, all = FALSE, plot
    = FALSE)
)
```
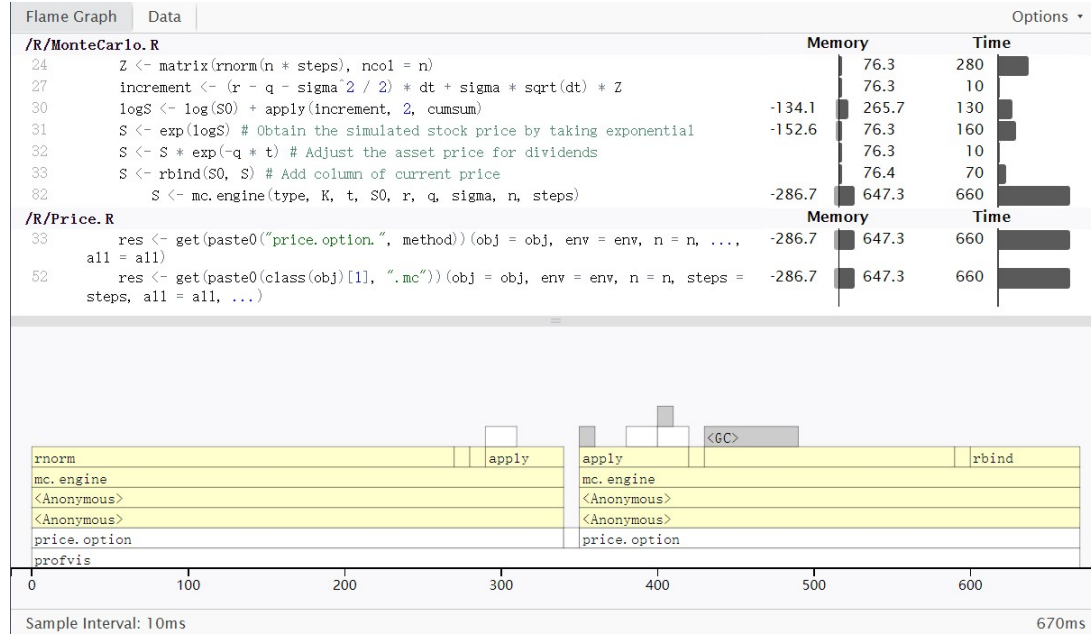
Figure 3.4: Profiling output of Monte Carlo pricing function

The named strips in Figure 3.4, represents different components that are triggered when running `price.option`. Each function represented by the corresponding strip is the internal component of the one below it. Here, we see `mc.engine` is the most (and only) time-consuming component of the pricing function. This means our prior adjustment of replacing for-loops with more efficient matrix operation is essential for reducing the computing cost of the whole pricing procedure.

An additional component is added to calculate the Standard Error of estimation. The formula used is shown by Equation 1.12:

```
1   mc.SE <- function(C_i, C.mean, n) {
2       s_c <- sqrt(sum(C_i - C.mean)^2 / (n - 1))
3       SE <- s_c / sqrt(n)
4       SE
5   }
```

The resulting standard error can be used to construct confidence intervals of the estimate of the option price.

One should be aware that Monte Carlo option pricing for American options is beyond the scope of this report. In the pricing functions below, we will only consider the pricing of European options.

3.7.2   vanilla.mc

The pricing of Vanilla options is exemplified in Equation 1.11. To obtain the simulated S_T's, one shall simply take the last row of the matrix output from the `mc.engine` and directly translate the formula to R code:

```
1   if (type == "call") {
2       C <- exp(-r * t) * pmax(S[steps + 1, ] - K, 0)
3   } else if (type == "put") {
4       C <- exp(-r * t) * pmax(K - S[steps + 1, ], 0)
5   }
```

29

Where the resulting discounted payoffs for each trajectory are stored in the vector `C`.

Notice that for a path-independent option such as a Vanilla option, we are only considering the option price at expiration. Thus, it is reasonable to take `step = 1` to reduce computational cost.

3.7.3 asian.mc

Exotic options such as Asian options and those included below are path-dependent. This implies capturing the price movements throughout time $t \in [0, T]$ is necessary.

Recall from Section 1.3, an Asian option's payoff is relevant to the underlying asset's average price through $t = [0, T]$. Thus, we will first compute and store the average prices of each trajectory.

```
1   S.mean <- apply(S, 2, mean)
```

The `apply` function with the second argument set to 2 allows us to apply the operation `mean` on each column of `S`, thus giving us the average price of each trajectory.

After obtaining the average, we will directly implement the payoff functions shown in List 1.3.1 in R as shown below:

```
1    if (is.avg_price) { # Compute payoff for average price Asian option
2        if (type == "call") {
3            C <- exp(-r * t) * pmax(S.mean - K, 0)
4        } else if (type == "put") {
5            C <- exp(-r * t) * pmax(K - S.mean, 0)
6        }
7    } else { # Compute payoff for average strike Asian option
8        if (type == "call") {
9            C <- exp(-r * t) * pmax(S[steps + 1, ] - S.mean, 0)
10       } else if (type == "put") {
11           C <- exp(-r * t) * pmax(S.mean - S[steps + 1, ], 0)
12       }
13   }
```

Notice that for path-dependent cases, using a large enough number of steps for each simulation to better resemble the real-world scenario is necessary to obtain an accurate approximation.

3.7.4 barrier.mc

As we see from List 1.3.1, a barrier option is "switched on and off" depending on whether the price movements crossed some pre-defined price barrier. To record this pre-specified quantity, we require users to include a property named `barrier` when they define their option object. We also require users to specify a boolean argument named `is.knockout` to classify whether the barrier option is "knock-in" or "knock-out". This is demonstrated by the below example:

```
1    option(style = "european", option = "barrier", type = "call", K = 20, t
     = 1, barrier = 21, is.knockout = TRUE)
```

Here, besides the regular parameters to see in any option, the user-defined the barrier price to be 21 and the option to be a "knock-out" barrier option, meaning that the option payoff will be zero if the price movement crossed the value 21.

To determine whether a price trajectory crossed a certain level, we can record the minimum and maximum through the trajectory and see if the barrier value lies between two extremas. If so,

because the price trajectory is continuous, we know that the trajectory must have crossed the barrier at least once through the option's lifetime.

We will use a similar `apply` method as mentioned in the previous section. Only now, we are computing the minimums and maximums of the trajectories. We then use a boolean vector `is.cross` to determine whether the barrier is bounded by the extremas.

```
1   S.min <- apply(S, 2, min)
2   S.max <- apply(S, 2, max)
3   is.cross <- (S.min < barrier & S.max > barrier)
```

We know a barrier option would behave like a vanilla option if ignoring the "knock-in/out" effect, so we will first calculate its discounted payoff as if dealing with a vanilla option.

```
1   if (type == "call") {
2       C <- exp(-r * t) * pmax(S[steps + 1, ] - K, 0)
3   } else if (type == "put") {
4       C <- exp(-r * t) * pmax(K - S[steps + 1, ], 0)
5   }
```

Then, we take into account the special effect and decide whether to zero the option payoff.

```
1   if (is.knockout) { # Case for knock-out barrier option
2       C <- ifelse(is.cross, 0, C)
3   } else { # Case for knock-in barrier option
4       C <- ifelse(is.cross, C, 0)
5   }
```

The function `ifelse(condition, A, B)` is simply a quick command for the regular if-else statement. It executes expression `A` if `condition` is evaluated to TRUE, and execute `B` otherwise.

3.7.5   binary.mc

As we see from List 1.3.1, a binary option, unlike other options discussed in the report, has a fixed payoff. To record this characteristic, we require users to pass in a value for argument `payout` when defining a binary option object, as shown in the example below:

```
1   option(style = "european", option = "binary", type = "call", K = 20, t =
    1, payout = 5)
```

For each trajectory, to decide whether the payoff should be zero or the specified value, we simply compare the simulated price at maturity S_T with the strike price K. This characteristic also makes a regular binary option path-independent.

Since the payoff is fixed by the variable `payout`, for simplicity, we can first set the discounted payoff of every trajectory to be equal, and reduced the "turned-off" cases to zero afterwards:

```
1   # Genearte a column vector with the same discounted payoff
2   C <- matrix(exp(-r * t) * payout, nrow = n, ncol = 1)
3
4   # Based on the specific type of the option, reduced corresponding rows
    to zero based on the payoff formula for binary option
5   if (type == "call") {
```

```
6         C <- ifelse(S[steps + 1, ] > K, C, 0)
7    } else if (type == "put") {
8         C <- ifelse(S[steps + 1, ] < K, C, 0)
9    }
```

3.7.6   lookback.mc

A lookback option is another case of a path-dependent option. As we see from List 1.3.1, a fixed strike lookback's payoff is determined by comparing the price extremas and the strike price throughout the lifetime, and a floating strike lookback's payoff is by comparing the price extremas and the price at maturity. This makes the lookback option path-dependent.

To determine whether the pre-specified lookback option is fixed or floating, we require the user to pass in a value for the boolean variable `is.fixed`, which indicates the option is "fixed strike" if evaluated to TRUE, and being "floating strike" otherwise. This is shown in the codes below:

```
1    option(style = "european", option = "lookback", type = "call", K = 20, t
     = 1, is.fixed = TRUE)
```

To allow us to calculate the payoff for each trajectory, we will first store the required value of minimum and maximum prices for each simulation, this could be done using the similar `apply` method as demonstrated above:

```
1    S.min <- apply(S, 2, min)
2    S.max <- apply(S, 2, max)
```

And the vector containing the simulated prices at maturity can be done with the same `S[steps + 1, ]` command.

After obtaining these elements, based on the classification of lookback options, there are 4 scenarios for the payoff to be calculated. By directly implementing the formulas as displayed in 1.3.1, we will obtain the vector `C` of discounted payoffs.

```
1    if (type == "call") {
2        if (is.fixed) { # Fixed strike
3            C <- exp(-r * t) * pmax(S.max - K, 0)
4        } else { # Float strike
5            C <- exp(-r * t) * (S[steps + 1, ] - S.min)
6        }
7    } else if (type == "put") {
8        if (is.fixed) { # Fixed strike
9            C <- exp(-r * t) * pmax(K - S.min, 0)
10       } else { # Float strike
11           C <- exp(-r * t) * (S.max - S[steps + 1, ])
12       }
13   }
```

Chapter 4

Results

This Chapter will discuss the outcome of the programming implementation discussed in previous chapters. The "outcome" will be presented using numerous option pricing examples based on our package, as well as some comments on the effectiveness of the result.

The second part of the chapter will discuss the limitations of the package and potential directions for future development.

4.1 Examples

This section gives a demonstration of options pricing via `rmcop` package. The installation of the package can be found in the README.md file in Appendix A.1. An alternative source would be the GitHub repository page which can be found here.

The following examples assumed users have `rmcop` installed and attached. Their device and software configuration should also be suitable (see Section 3.2 for reference) to run the package.

Before running any of the following programmes, users shall ensure that they have attached `rmcop` to their current environment, such as using the command:

```
1   library(rmcop)
```

4.1.1 Defining Option and Environment Objects

As mentioned in Section 3.3, `"option"` and `"env"` class objects are defined using functions `option` and `option.env`, respectively.

For example, if we want to define an *knock-out Barrier European call option* with a strike price of 20, maturity of 1 year, and barrier of 10, we can use the following command:

```
1   option(
2       # Common arguments for the father class "option"
3       style = "european", type = "call", option = "barrier", K = 20, t =
        1,
4       # Arguments for the subclass "barrier option"
5       barrier = 10, is.knockout = TRUE
6   )
```

And the result S3 list would be:

```
1   > $style
2   > [1] "european"
3   >
4   > $type
5   > [1] "call"
6   >
7   > $K
8   > [1] 20
9   >
10  > $t
11  > [1] 1
12  >
13  > $barrier
14  > [1] 10
15  >
16  > $is.knockout
17  > [1] TRUE
18  >
19  > attr(,"class")
20  > [1] "barrier" "option"
```

If we wish to define an option of other types, styles, or families, simply adjusted the "common" arguments and specify additional subclass-specific arguments for exotic options.

To define an `env` class object, for example, a market environment where the current price is 30, the interest rate of 0.03, the dividend yield rate of 0.02, and the volatility rate of 0.01 can be defined using the code below:

```
1   option.env(S = 30, r = 0.03, q = 0.02, sigma = 0.01)
```

4.1.2   Black-Scholes

`rmcop` only supports Black-Scholes formula pricing for vanilla options[1].

Suppose we want to know the price of a vanilla option which is specified as:

```
1   option.vanilla <- option("european", "vanilla", "call", K = 20, t = 1)
```

To use the Black-Scholes method, we can specify the argument `method = "bs"` when defining the environment object or when triggering the pricing function.

```
1   # Specified when defining environment
2   env <- option.env(method = "bs", S = 15, r = 0.5, sigma = 0.1)
3
4   # Specified when calling pricing function
5   price.option(obj = option.vanilla, env = env, method = "bs")
6   > [1] 2.877561
```

The result price will then be returned after calling the pricing funtion.

---

[1]This includes vanilla European options and American non-dividend bearing call option, as other American options cannot be solved using explicit form

4.1.3   Binomial Tree

The package supports pricing both European and American vanilla options. Similar to the Black-Scholes case, to change the method to using binomial lattice tree pricing, simply specify the argument `method = "binomial"` when defining the environment object or calling the pricing function. We will use the same option-environment setup as in the previous section.

Because the binomial tree method used different arguments than the Black-Scholes setup used for the Black-Scholes method and the Monte Carlo method, when calling the pricing function, users need to specify a few more arguments, including ratios of upward and downward jumps.

```
1  price.option(obj = option.vanilla, env = env, method = "bs")
```

The below example uses the binomial tree method to price an option under the same setup.

```
1  price.option(
2      obj = option.vanilla, env = env, method = "binomial",
3      u = 1.2, # Ratio of upward jump
4      d = 0.9, # Ratio of downward jump
5      n = 10 # Number of steps
6  )
7  > [1] 4.166904
```

The Binomial tree method also allows users to extract more information during the calculation. This can be achieved by specifying the argument `all = TRUE` in the pricing function.

```
1   # Changed n = 4 for better demonstration
2   res <- price.option(obj = option.vanilla, env = env, method =
    "binomial", u = 1.2, d = 0.9, n = 4, all = TRUE)
3
4   res$p # Extract the risk-neutral probability
5   > [1] 0.7771615
6   res$price_tree # Extract the binomial price tree matrix
7   >           [,1]    [,2]    [,3]    [,4]    [,5]
8   > [1,] 15.0000     NA      NA      NA      NA
9   > [2,] 13.5000 18.000     NA      NA      NA
10  > [3,] 12.1500 16.200 21.600     NA      NA
11  > [4,] 10.9350 14.580 19.440 25.920     NA
12  > [5,]  9.8415 13.122 17.496 23.328 31.104
```

4.1.4   Monte Carlo

To demonstrate the pricing function under the Monte Carlo setup, suppose we are interested in a *binary put option* with a strike price of 50, maturity of 2 years and payout of 10, under a market environment where the current price is 15, interest rate 0.5, and volatility 0.1. We will encode this information using the code below:

```
1  op.binary <- option("european", "binary", "put", K = 50, t = 2, payout =
   10)
2  env <- option.env(method = "mc", S = 15, r = 0.5, sigma = 0.1)
```

We will price this option-environment combination using the Monte Carlo method under the

setup of 10000 trajectories simulations, and 1 step only for each trajectory, since a binary option is path-independent.

```
1  set.seed(10)
2  price.option(op.binary, env, n = 1000, steps = 1)
3  > [1] 3.487497
```

And we will obtain the estimated price.

We can also request for returning more information about the result. For example, we may wish to plot the price trajectories. We can achieve this my simply adding the argument `plot = TRUE`.

```
1  # For better demonstration, we set steps = 50
2  price.option(op.binary, env, n = 1000, steps = 50, plot = TRUE)
```

This will return us an estimate of the price as well as a result plot of the price trajectory, as shown by Figure 4.1.
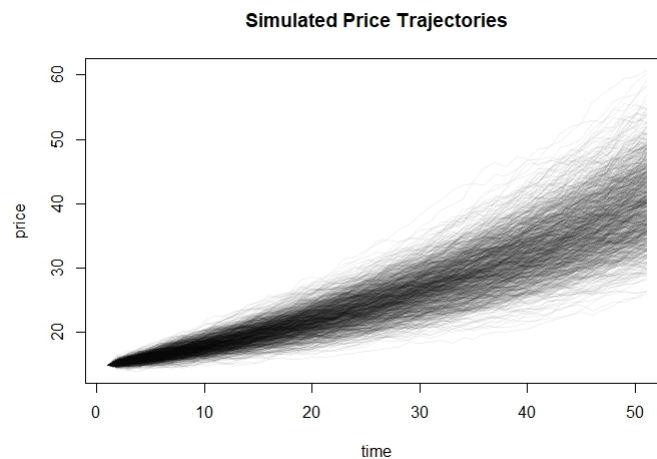

Figure 4.1: Plot of Trajectories

We can also use argument `all = TRUE` to store some other useful information (such as option's discounted payoffs) for each simulation[2].

```
1  res <- price.option(op.binary, env, n = 1000, steps = 1, all = TRUE)
2  head(res$C) # C is the keyword for extracting option payoffs
3  > [1] 3.678794 3.678794 3.678794 3.678794 3.678794 3.678794
4  summary(res$C)
5  >  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
6  > 0.000   3.679   3.679   3.447   3.679   3.679
```

4.1.5 Comparing Multiple Option-Environment Setups

One of the most useful properties of encapsulating options and environment characteristics using OOP is that we can compare different setups without adjusting too many arguments. In real

---

[2]We see from the first few items of the list that the discounted payoff is either identical or zero (zero cases happens to not be in the head of the vector), which is what we would expect to see from a binary option

option pricing scenarios, investors or hedgers[3] may be interested in comparing multiple financial options on the same underlying asset.

For example, if we are to compare three different options' prices given the same market condition (i.e. price condition of the underlying asset), shown by the codes below: We can also use argument `all = TRUE` to store some other useful information (such as option's discounted payoffs) for each simulation[4].

```
# Defining options of interest
op1 <- option("european", "vanilla", "call", 10, 2)
op2 <- option("european", "vanilla", "put", 12, 2)
op3 <- option("european", "binary", "call", 10, 2, payout = 5)
op.list <- list(op1, op2, op3) # Create a list contining the three
    options of interest

# Defining the fixed environment
env <- option.env(method = "mc", S = 10, r = 0.05, sigma = 0.01, n =
    100000, steps = 1)
```

We can obtain the result of the three options' corresponding prices by using the `lapply` method with just one line of code:

```
set.seed(10) # Set seed for random number generator
unlist(lapply(op.list, price.option, env = env))
> [1] 0.9518262 0.8611994 4.5241871

# Equivalently, we can write the pricing function separatedly for each
    option
price.option(op1, env)
price.option(op2, env)
price.option(op3, env)
```

We can apply similar methods for cases when comparing different market environment setups. As long as the objects are well defined, one can price any option-environment combinations by simply passing in just two arguments for the `price.option` functions.

4.2 Limitations and Future Development

Currently, `rmcop` only supports a few types of options. The derivative market is varying, and it is infeasible for a single package to contain tools that can address all possible scenarios. One of the necessary future development of `rmcop` will be extending its support to more cases. This includes the support of pricing American and Bermudian styles options via the Monte Carlo method, and the support of more types of exotic options pricing.

Additional techniques enhancing the efficiency of Monte Carlo option pricing, including variance reduction techniques or quasi-Monte Carlo methods [2] are not yet implemented by this package. These features may be included in future versions of the package.

Though the development of the package includes several R-programming tricks to speed up the computation, one cannot guarantee the implementation here is the most efficient way of encoding the pricing methods. Also, as we discussed in Section 2, `rmcop` being an R-based package has fundamentally no computational advantage compared to algorithms developed via lower-level

---

[3]The party who wishes to use options to offset the risk against market fluctuation.

[4]We see from the first few items of the list that the discounted payoff is either identical or zero (zero cases happens to not be in the head of the vector), which is what we would expect to see from a binary option

programming languages. The main focus of this package is for the demonstration of model implementations, and the C++ programmes should be more desirable when conducting more practical and computationally-intensive option pricing projects.

Until the time when this report is completed, the development of `rmcop` is still active. One may find resources in Appendix A useful in terms of tracking future updates of the package.

Appendix A

Appendix

A.1 Package Source Code

This appendix includes *selected* key components of the package functions. Notice that for code readability in this PDF file, some functions and function comments are omitted. Therefore, visiting the package's GitHub Repository is highly recommended. The online repository stores the full directory of the package with detailed installation instructions. This, however, comes with the risk of breaking anonymity, as some parts of the GitHub site may reveal the creator's information.

A.2 Option.R

```r
1   # Creating new option object
    ====================================================
2   #' Initialise new `option` class object
3   #' Create an S3 class object which characterised an option of interest.
    The returned object is then used together with a `"env"` class object
    together as inputs of the `price.option` function for further option
    pricing.
4   option <- function(style, option = "vanilla", type, K, t, ...) {
5
6       obj <- list(
7           "style" = style,
8           "type" = type,
9           "K" = K,
10          "t" = t
11      )
12
13      # Distribute to corresponding initialisation function according to
        specified arguments
14      if (option == "vanilla") {
15          obj <- option.vanilla(obj)
16      } else {
17          obj <- option.exotic(obj, option, ...)
18      }
19      obj
20  }
21
22  ## Vanilla
    =======================================================================
23  #' Initialise vanilla option
24  option.vanilla <- function(obj) {
25      class(obj) <- c("vanilla", "option")
26      obj
27  }
28
29  ## Exotic
    =======================================================================
30  #' Initialise exotic option
31  option.exotic <- function(obj, option, ...) {
32      obj <- get(paste0("option.exotic.", option))(obj, ...) # Direct to
        corresponding initialisation function
33      obj
34  }
35
36  #' Initialise Asian option
37  option.exotic.asian <- function(obj, is.avg_price = TRUE) {
38      obj <- c(
39          obj,
40          "is.avg_price" = is.avg_price
41      )
42      class(obj) <- c("asian", "option")
43      obj
44  }
```

```r
#' Initialise barrier option
option.exotic.barrier <- function(obj, barrier, is.knockout = TRUE) {
    obj <- c(
        obj,
        "barrier" = barrier,
        "is.knockout" = is.knockout
    )
    class(obj) <- c("barrier", "option")
    obj
}

#' Initialise binary option
option.exotic.binary <- function(obj, payout) {
    obj <- c(
        obj,
        "payout" = payout
    )
    class(obj) <- c("binary", "option")
    obj
}

#' Initialise lookback option
option.exotic.lookback <- function(obj, is.fixed = TRUE) {
    obj <- c(
        obj,
        "is.fixed" = is.fixed
    )
    class(obj) <- c("lookback", "option")
    obj
}

# Creating new option environment
===============================================

#' Initialise new `env` class object
#' Create an S3 class object which characterised the market environment
used when pricing the option of interest. The returned object is used
together with an `"option"` class object as inputs of the pricing
function `price.option`.
option.env <- function(method = "mc", S, r, q = 0, sigma, n = NULL,
steps = NULL) {
    env <- list(
        "method" = method,
        "S" = S,
        "r" = r,
        "q" = q,
        "sigma" = sigma,
        "n" = n,
        "steps" = steps
    )
    class(env) <- "env"
    env
}
```

A.3 Price.R

```r
1   # Top Level Pricing Functions for Options
    =======================================
2
3   #' Pricing `option` Class Object
4   #' The all-in-one pricing function for `rmcop`. User should predefine an
    `"option"` class object and an `"env"` class object as inputs of
    `price.option`. According to the pricing method selected, the user shall
    provide additional inputs as instructions for the pricing function.
5   price.option <- function(obj, env, method = env$method, n = env$n, ... ,
    all = FALSE) {
6       check.method(method)
7       res <- get(paste0("price.option.", method))(obj = obj, env = env, n
        = n, ..., all = all)
8       res
9   }
10
11  #' Option Pricing via Monte Carlo Method
12  price.option.mc <- function(obj, env, n, steps = env$steps, all, ...) {
13      if (is.null(n)) stop("n is not specified")
14      if (is.null(steps)) {
15          steps <- 1
16          warning("steps is not specified, evaluated to default value 1")
17      }
18      res <- get(paste0(class(obj)[1], ".mc"))(obj = obj, env = env, n =
        n, steps = steps, all = all, ...)
19      res
20  }
21
22  #' Option Pricing via Black-Scholes Model
23  price.option.bs <- function(obj, env, n, all) {
24      if (obj$style == "american") {
25          if (env$q != 0 | obj$type != "call") {
26              warning("Black Scholes valuation should only apply on
                  European european options or American call option with
                  non-dividend paying asset, the result here will be based on
                  European options pricing and may not be accurate.")
27          }
28      }
29      res <- get(paste0(class(obj)[1], ".bs"))(obj, env, all)
30      res
31  }
32
33  #' Option Pricing via Binomial Lattice Model
34  price.option.binomial <- function(obj, env, n, u, d, all) {
35      res <- get(paste0(class(obj)[1], ".binomial"))(obj = obj, env = env,
        n = n, u = u, d = d, all = all)
36      res
37  }
```

A.4   MonteCarlo.R

```r
#' Engine for Monte Carlo Simulation on Asset Price Trajectory
mc.engine <- function(type, K, t, S0, r, q, sigma, n, steps) {

    dt <- t / steps

    # Generate n random samples from N(0,1)
    Z <- matrix(rnorm(n * steps), ncol = n)

    # Get logarithm of price change per step
    increment <- (r - q - sigma^2 / 2) * dt + sigma * sqrt(dt) * Z

    # Use vectorised MC method to compute logarithm of S(t) at each time
    step
    logS <- log(S0) + apply(increment, 2, cumsum)
    S <- exp(logS) # Obtain the simulated stock price by taking
    exponential
    S <- S * exp(-q * t) # Adjust the asset price for dividends
    S <- rbind(S0, S) # Add column of current price
    S
}

#' Compute Monte Carlo Model's Standard Error
mc.SE <- function(C_i, C.mean, n) {
    s_c <- sqrt(sum(C_i - C.mean)^2 / (n - 1))
    SE <- s_c / sqrt(n)
    SE
}

# Model
=========================================================================

#' Pricing vanilla option using Monte Carlo method
vanilla.mc <- function(obj, env, n, steps, all = FALSE, plot = FALSE,
...) {

    # Check object integrity, CAN BE OMITTED
    check.args(obj, env, option = "vanilla")

    # Extract objects properties
    style <- obj$style
    type <- obj$type
    K <- obj$K
    t <- obj$t
    S0 <- env$S
    r <- env$r
    q <- env$q
    sigma <- env$sigma

    # Apply MC
    if (style == "european") {
        S <- mc.engine(type, K, t, S0, r, q, sigma, n, steps)
    } else if (style == "american") {
```

```r
            stop("American option pricing via MC not supported")
        }

        # Compute corresponding discounted price at current time
        # Vanilla option only consider option price at maturity, so only
        #   take the last column of the price matrix S_i
        if (type == "call") {
            C <- exp(-r * t) * pmax(S[steps + 1, ] - K, 0)
        } else if (type == "put") {
            C <- exp(-r * t) * pmax(K - S[steps + 1, ], 0)
        }

        # Obtain the estimate of option price by taking average of C
        C.mean <- mean(C, ...)

        # Compute standard error of the estimate
        SE <- mc.SE(C, C.mean, n)

        # Plot function
        if (plot) {mc.plot(S, ...)}

        # Return result
        if (all) {
            # Return actual object environment used
            env <- option.env(method = "mc", S = S, r = r, q = q, sigma =
            sigma, n = n, steps = steps)
            list(
                "env" = env,
                "price" = C.mean,
                "SE" = SE,
                "S" = S,
                "C" = C
            )
        } else {
            C.mean
        }
    }

#' Pricing Asian option using Monte Carlo method
asian.mc <- function(obj, env, n, steps, all = FALSE, plot = FALSE, ...)
{
        # Check object integrity, CAN BE OMITTED
        check.args(obj, env, option = "asian")

        # Extract objects properties
        style <- obj$style
        type <- obj$type
        K <- obj$K
        t <- obj$t
        is.avg_price <- obj$is.avg_price
        S0 <- env$S
        r <- env$r
        q <- env$q
        sigma <- env$sigma
```

```r
        # Apply MC
        if (style == "european") {
            S <- mc.engine(type, K, t, S0, r, q, sigma, n, steps)
        } else if (style == "american") {
            stop("American option pricing via MC not supported")
        }

        # Compute the mean of each replication's price movement, as required
        to compute Asian option pay-off
        S.mean <- apply(S, 2, mean)

        # Compute corresponding discounted price at current time
        if (is.avg_price) { # Compute payoff for average price Asian option
            if (type == "call") {
                C <- exp(-r * t) * pmax(S.mean - K, 0)
            } else if (type == "put") {
                C <- exp(-r * t) * pmax(K - S.mean, 0)
            }
        } else { # Compute payoff for average strike Asian option
            if (type == "call") {
                C <- exp(-r * t) * pmax(S[steps + 1, ] - S.mean, 0)
            } else if (type == "put") {
                C <- exp(-r * t) * pmax(S.mean - S[steps + 1, ], 0)
            }
        }

        # Obtain the estimate of option price by taking average of C
        C.mean <- mean(C)

        # Compute standard error of the estimate
        SE <- mc.SE(C, C.mean, n)

        # Plot function
        if (plot) {mc.plot(S, ...)}

        # Return result
        if (all) {
            # Return actual object environment used
            env <- option.env(method = "mc", S = S, r = r, q = q, sigma =
            sigma, n = n, steps = steps)
            list(
                "env" = env,
                "price" = C.mean,
                "SE" = SE,
                "S" = S,
                "C" = C
            )
        } else {
            C.mean
        }
    }


    #' Pricing Barrier option using Monte Carlo method
    barrier.mc <- function(obj, env, n, steps, all = FALSE, plot = FALSE,
    ...) {
```

```r
153         # Check object integrity, CAN BE OMITTED
154         check.args(obj, env, option = "barrier")
155
156         # Extract objects properties
157         style <- obj$style
158         type <- obj$type
159         K <- obj$K
160         t <- obj$t
161         barrier <- obj$barrier
162         is.knockout <- obj$is.knockout
163         S0 <- env$S
164         r <- env$r
165         q <- env$q
166         sigma <- env$sigma
167
168         # Apply MC
169         if (style == "european") {
170             S <- mc.engine(type, K, t, S0, r, q, sigma, n, steps)
171         } else if (style == "american") {
172             stop("American option pricing via MC not supported")
173         }
174
175         # Check if price movements corssed the specified barrier
176         S.min <- apply(S, 2, min)
177         S.max <- apply(S, 2, max)
178         is.cross <- (S.min < barrier & S.max > barrier)
179
180         # Compute corresponding discounted price at current time
181         if (type == "call") {
182             C <- exp(-r * t) * pmax(S[steps + 1, ] - K, 0)
183         } else if (type == "put") {
184             C <- exp(-r * t) * pmax(K - S[steps + 1, ], 0)
185         }
186
187         # Adjust for knock-out & knock-in situations for a barrier option
188         if (is.knockout) {
189             C <- ifelse(is.cross, 0, C)
190         } else {
191             C <- ifelse(is.cross, C, 0)
192         }
193
194         # Obtain the estimate of option price by taking average of C
195         C.mean <- mean(C)
196
197         # Compute standard error of the estimate
198         SE <- mc.SE(C, C.mean, n)
199
200         # Plot function
201         if (plot) {mc.plot(S, ...)}
202
203         # Return result
204         if (all) {
205             # Return actual object environment used
206             env <- option.env(method = "mc", S = S, r = r, q = q, sigma =
                sigma, n = n, steps = steps)
```

```r
          list(
              "env" = env,
              "price" = C.mean,
              "SE" = SE,
              "S" = S,
              "C" = C
          )
      } else {
          C.mean
      }
  }

  #' Pricing Barrier option using Monte Carlo method
  binary.mc <- function(obj, env, n, steps, all = FALSE, plot = FALSE,
  ...) {

      # Check object integrity, CAN BE OMITTED
      check.args(obj, env, option = "binary")

      # Extract objects properties
      style <- obj$style
      type <- obj$type
      K <- obj$K
      t <- obj$t
      payout <- obj$payout
      S0 <- env$S
      r <- env$r
      q <- env$q
      sigma <- env$sigma

      # Apply MC
      if (style == "european") {
          S <- mc.engine(type, K, t, S0, r, q, sigma, n, steps)
      } else if (style == "american") {
          stop("American option pricing via MC not supported")
      }

      # Compute corresponding discounted price at current time
      # Compute uniform payout for each MC replication of a binary option
      C <- matrix(exp(-r * t) * payout, nrow = n, ncol = 1)
      if (type == "call") {
          C <- ifelse(S[steps + 1, ] > K, C, 0)
      } else if (type == "put") {
          C <- ifelse(S[steps + 1, ] < K, C, 0)
      }

      # Obtain the estimate of option price by taking average of C
      C.mean <- mean(C)

      # Compute standard error of the estimate
      SE <- mc.SE(C, C.mean, n)

      # Plot function
      if (plot) {mc.plot(S, ...)}
```

```r
        # Return result
        if (all) {
            # Return actual object environment used
            env <- option.env(method = "mc", S = S, r = r, q = q, sigma =
            sigma, n = n, steps = steps)
            list(
                "env" = env,
                "price" = C.mean,
                "SE" = SE,
                "S" = S,
                "C" = C
            )
        } else {
            C.mean
        }
    }

    #' Pricing Lookback option using Monte Carlo method
    lookback.mc <- function(obj, env, n, steps, all = FALSE, plot = FALSE,
    ...) {
        # Check object integrity, CAN BE OMITTED
        check.args(obj, env, option = "lookback")

        # Extract objects properties
        style <- obj$style
        type <- obj$type
        K <- obj$K
        t <- obj$t
        is.fixed <- obj$is.fixed
        S0 <- env$S
        r <- env$r
        q <- env$q
        sigma <- env$sigma

        # Apply MC
        if (style == "european") {
            S <- mc.engine(type, K, t, S0, r, q, sigma, n, steps)
        } else if (style == "american") {
            stop("American option pricing via MC not supported")
        }

        # Check if price movements crossed the specified barrier
        S.min <- apply(S, 2, min)
        S.max <- apply(S, 2, max)

        # Compute corresponding discounted price at current time
        if (type == "call") {
            if (is.fixed) {
                C <- exp(-r * t) * pmax(S.max - K, 0)
            } else {
                C <- exp(-r * t) * (S[steps + 1, ] - S.min)
            }
        } else if (type == "put") {
            if (is.fixed) {
                C <- exp(-r * t) * pmax(K - S.min, 0)
```

```
314        } else {
315            C <- exp(-r * t) * (S.max - S[steps + 1, ])
316        }
317    }
318
319    # Obtain the estimate of option price by taking average of C
320    C.mean <- mean(C)
321
322    # Compute standard error of the estimate
323    SE <- mc.SE(C, C.mean, n)
324
325    # Plot function
326    if (plot) {mc.plot(S, ...)}
327
328    # Return result
329    if (all) {
330        # Return actual object environment used
331        obj <- option(style = style, option = "lookback", type = type, K
               = K, t = t, is.fixed = is.fixed)
332        env <- option.env(method = "mc", S = S, r = r, q = q, sigma =
               sigma, n = n, steps = steps)
333        list(
334            "env" = env,
335            "C.price" = C.mean,
336            "SE" = SE,
337            "S" = S,
338            "C" = C
339        )
340    } else {
341        C.mean
342    }
343 }
```

A.5   BlackScholes.R

```r
#' Compute d1 in Black-Scholes Formula
bs.d1 <- function(S, K, t, r, q, sigma) {
    (log(S / K) + (r - q + sigma^2 / 2) * t) / (sigma * sqrt(t))
}

# Model
=======================================================================

#' Pricing vanilla option using Black-Scholes model
vanilla.bs <- function(obj, env, all = FALSE) {
    type <- obj$type
    K <- obj$K
    t <- obj$t
    S <- env$S
    r <- env$r
    q <- env$q
    sigma <- env$sigma

    d1 <- bs.d1(S, K, t, r, q, sigma)
    d2 <- d1 - sigma * sqrt(t)

    if (type == "call") {
        price <- exp(-q * t) * S * pnorm(d1) - exp(-r * t) * K *
        pnorm(d2)
    } else if (type == "put") {
        price <- exp(-r * t) * K * pnorm(-d2) - exp(-q * t) * S *
        pnorm(-d1)
    }

    if (all) {
        list(
            obj,
            env,
            "price" = price
        )
    } else {
        price
    }
}
```

A.6 Binomial.R

```r
# Engine
========================================================================
#' Generating Binomial Tree Matrix
Binomial.tree <- function(S0, u, d, n) {
    if (is.null(n)) stop("Number of binomial tree steps n is not
    defined.")
    dim <- n + 1
    S <- matrix(nrow = dim, ncol = dim) # Predefine the size
    temp <- c(S0)
    for (i in 1:dim) {
        S[i, ] <- c(temp, rep(NA, dim - i))
        temp <- c(temp[1] * d, temp * u)
    }
    S
}

#' Backward Recursive Risk-neutral Valuation for Binomial Model
Binomial.rnv <- function(fn, p, q, r, n, dt, K = NULL, S = NULL,
                         type = "call",
                         style = "european") {
    if (style == "european") {
        rnv.european <- function(fn, p, r, dt, state = n) {
            if (state <= 0) {
                fn
            } else {
                l <- state # Specify the length l of the vector
                containing prices of next state looking back by 1 step
                fm <- sapply(1:l, function(i) {exp(-r * dt) * (p * fn[i
                + 1] + q * fn[i])})
                rnv.european(fm, p, r, dt, state = state - 1)
            }
        }
        rnv.european(fn, p, r, dt)
    } else if (style == "american") {
        rnv.american <- function(fn, p, r, dt, K, S, type = type, state
        = n) {
            if (state <= 0) {
                fn
            } else {
                l <- state
                fm <- sapply(1:l, function(i) {
                    if (type == "call") {
                        max(max(S[state, i] - K, 0), exp(-r * dt) * (p *
                        fn[i + 1] + q * fn[i]))
                    } else if (type == "put") {
                        max(max(K - S[state, i], 0), exp(-r * dt) * (p *
                        fn[i + 1] + q * fn[i]))
                    }
                })
                rnv.american(fm, p, r, dt, K, S, type = type, state =
                state - 1)
            }
```

```r
        }
        rnv.american(fn, p, r, dt, K, S, type = type)
    } else {
        stop("Invalid option style. `style` must be specified to either
            'european' or 'american'.")
    }
}

# Model
==============================================================================

#' Binomial model for pricing vanilla options
vanilla.binomial <- function(obj, env, n, u, d, all = FALSE) {

    style = obj$style
    type = obj$type
    K = obj$K
    t = obj$t
    S = env$S
    r = env$r

    dt <- t / n # Compute delta t
    p <- (exp(r * dt) - d) / (u - d) # Compute riskless probability for
    an upward price movement
    q <- 1 - p # Compute riskless probability for a downward price
    movement

    # Store the stock prices of the tree for each state in a matrix
    `price_tree`
    price_tree <- Binomial.tree(S, u, d, n)
    Sn <- price_tree[n + 1, ] # Stores the possible stock prices at the
    terminal steps.

    # Calculate the option value at the final step, i.e. maturity.
    if (type == "call") {
        fn <- sapply(Sn, function(x) {max(x - K, 0)})
    } else if (type == "put") {
        fn <- sapply(Sn, function(x) {max(K - x, 0)})
    } else {
        stop("Invalid option type. `type` must be specified to either
            'call' or 'put'.")
    }

    # Calculate the option value at current time.
    price <- Binomial.rnv(fn, p, q, r, n, dt, K, S = price_tree, type =
    type, style = style)

    # Return output
    if (all) {
        list(
            "price" = price,
            "p" = p,
            "price_tree" = price_tree
        )
    } else {
```

```
92              price
93          }
94      }
```

Bibliography

[1] D. Higham, *An Introduction to Financial Option Valuation*. Cambridge University Press, 2004. Exotic Options Pricing.

[2] P. Glasserman, *Monte Carlo methods in financial engineering*. Springer, 2003. Explaination on Monte Carlo method, MC European vanilla option pricing, MC Asian option pricing.

[3] S. M. Iacus, *Stochastic Processes and Stochastic Differential Equations*. Springer New York, 2008. Stochastic Processes¡br/¿Brownian Motion.

[4] S. E. Shreve, *Stochastic calculus for finance. II, Continuous-time models*. Springer, 2004. Geometric Brownian Motion¡br/¿Ito Calculus.

[5] L. Bachelier, "Théorie de la spéculation," *Annales scientifiques de l'École normale supérieure*, vol. 17, pp. 21–86, 1900. Earliest attempt of quantitative method in option pricing.

[6] C. M. Sprenkle, "Warrant prices as indicators of expectations and preferences," *Yale economic essays*, vol. 1, pp. 179–232, 1961. Sprenkle Formula.

[7] F. Black and M. Scholes, "The pricing of options and corporate liabilities," *The Journal of political economy*, vol. 81, pp. 637–654, 1973. Definition of Financial Options¡br/¿Black-Scholes Model.

[8] J. C. Cox, S. A. Ross, and M. Rubinstein, "Option pricing: A simplified approach," *Journal of financial economics*, vol. 7, pp. 229–263, 1979. Binomial Lattice Model.

[9] R. McDonald, *Package 'derivmkts'*. CRAN, 0.2.5 ed., 10 2022.

[10] O. Mersmann, C. Beleites, R. Hurling, A. Friedman, and J. M. Ulrich, *Package 'microbenchmark'*. CRAN, 1.4.9 ed., 10 2022.

[11] D. Wuertz, *Package 'fOptions'*. CRAN, 11 2017.

[12] D. Eddelbuettel, K. Nguyen, and T. Leitch, *Package 'RQuantLib'*. CRAN, 10 2022.

[13] H. Wickham, J. Hester, W. Chang, J. Bryan, and RStudio, *Package 'devtools'*. CRAN, 2.4.5 ed., 10 2022.

[14] H. Wickham, RStudio, and R. team, *Package 'testthat'*. CRAN, 3.1.7 ed., 3 2023.

[15] W. Chang, J. Luraschi, T. Mastny, RStudio, jQuery Foundation, jQuery contributors, M. Bostock, D. contributors, and I. Sagalaev, *Package 'profvis'*. CRAN, 0.3.7 ed., 10 2022.